

# Report: Fluid Simulation in Computer Graphics WS2020

## Weakly Compressible Smoothed Particle Hydrodynamics and Position Based Fluids

Berat Ertural, Dennis Ledwon  
RWTH Aachen, Visual Computing Institute  
Supervisor: Jose Antonio Fernandez Fernandez

Reading maketh a full man; conference a  
ready man; and writing an exact man.

*Sir Francis Bacon*

## 1 Introduction

Fluids play a vital role in our lives. Understanding the motion and behaviour of fluids is not only crucial to many fields of engineering, but it also aids us in creating beautiful virtual worlds, such as used in games and cinema. The simulation of fluids is a vast and rapidly growing field. During the course of this practical, we studied various methods, worked on simulations of our own and analysed different aspects a fluid solver.

The motion of fluids is governed by the famous Navier-Stokes equation, as given below in its Lagrangian form;

$$\dot{v}_i = \frac{\partial v_i}{\partial t} = -\frac{1}{\rho_i} \nabla p_i + \nu \nabla^2 v_i + \frac{F_i^{ext}}{m_i}. \quad (1)$$

The Lagrangian form denotes the sampling points as dynamic particles moving with the flow of the fluid. The above PDE describes the acceleration  $\dot{v}_i$  of such a particle  $i$  as depending on the internal pressure gradient  $\nabla p_i$  and the velocity divergence  $\nabla^2 v_i$  of the fluid. The velocity divergence and friction coefficient  $\nu$  represent the viscosity of the fluid due to internal friction. The pressure gradient drives the acceleration due to internal pressure differences, to achieve incompressibility. In theory the density of the fluid is said to be constant,  $\frac{\partial \rho}{\partial t} = 0$ . In practice, absolute incompressibility is not achievable due to numerical errors. Sometimes small deviations from the rest density of the fluid is desirable. This will be discussed in the next chapters. [BK15, IOS<sup>+</sup>14]

As with many PDE describing natural phenomena, the Navier-Stokes equation has no analytical solution. If it had, this report would end here with a reference to the proof of someone who is now one million dollars richer. Nonetheless, numerical approximations to fluid dynamics exist and during the course of the practical lecture we were tasked to implement two such solvers, namely Weakly Compressible Smoothed Particle Hydrodynamics (WCSPH) [IOS<sup>+</sup>14] and Position based Fluids (PBF) [MM13]. In the next section we will give a short overview of the used methods and implementation of the application. We will then present various experiments and discuss our observations and the results.

## 2 Methods

In the following a short overview of the methodological background and implementation will be given. In addition, certain aspects of the implementation that were done differently than from the given assignment will be presented. For a thorough discussion of the SPH and PBF framework we refer the reader to the review of SPH fluids by Ihmsen et al. [IOS<sup>+</sup>14] and to the original paper about PBF by Macklin et al. [MM13], as well as the works of Akinici et al. [AIA<sup>+</sup>12, AAT13].

### 2.1 Weakly Compressible SPH

Smoothed Particle Hydrodynamics (SPH) is a class of particle based methods for fluid simulation with a long record of scientific publications. The core idea of SPH is to approximate any attribute  $A_i$  of a point at an arbitrary position  $\vec{x}_i$  by weighting the attributes of the immediate sur-

rounding particles  $j \in N$ ;

$$A_i = \sum_{j \in N} \frac{m_j}{\rho_j} A_j \cdot W(\|\vec{x}_i - \vec{x}_j\|/h). \quad (2)$$

Particle mass  $m_j$  and density  $p_j$  is used as a normalization factor. The weighting is performed by the so called kernel function  $W : \mathbb{R} \rightarrow \mathbb{R}$ . It depends on the ratio of the distance between two particles and the smoothing length  $h$ , which is defined to be a multiple of the particle diameter;  $h = \mu \cdot 2r$ . We set  $\mu = 1.1$ . during the remainder of this report. The kernel is then given as  $W(q) = \frac{1}{h^3} f(q)$ . The function for the kernel used throughout the application is the following cubic spline [Mon92],

$$f(q) = \frac{3}{2\pi} \begin{cases} \frac{2}{3} - q^2 + \frac{1}{2}q^3 & 0 \leq q < 1 \\ \frac{1}{6}(2 - q)^3 & 1 \leq q < 2 \\ 0 & 2 \leq q \end{cases} \quad (3)$$

It should be noted, that the effective radius of the kernel function, and therefore the number of neighbourhood particles taken into account, depends on the smoothing length. Weakly Compressible SPH first computes the particle density  $\rho_i$  using (2) and then applies this to a so called Equation of State (EOS) to derive the resulting particle pressures  $p_i$ . In the application the following EOS was used to derive the pressures;

$$p_i = \max(0, B(\rho_i - \rho_0)), \quad (4)$$

where  $B$  is a stiffness constant. Equation (4) is similar to hookes law, as the pressure and therefore the resulting force depends linearly on the displacement. Therefore WCSPH allows for a small degree of incompressibility of the fluid. We will later discuss the implications of this. Using the formulations above we can derive expressions for the gradient and divergence field of an attribute and therefore compute the pressure and viscosity term of the Navier-Stokes equation.

## 2.2 Position Based Fluids

Position Based Fluids (PBF) is a relatively new method based on the authors previous work on Position Based Dynamics (PBD) [MHHR07].

Within this framework, a system is modelled using constraint functions, which describe the theoretical optimum or the equilibrium state. A solver will try to move the system towards the equilibrium by iteratively moving the positions of the sample points, such that the constraint violation is minimized. For fluids simulations, PBF tries to achieve incompressibility by specifying the constraint (for a particle  $i$ );

$$C_i(\hat{x}) = \frac{\rho_i(\hat{x})}{\rho_0} - 1, \quad (5)$$

where  $\hat{x}$  is the collection of all particles involved in the computation of  $\rho_i$ . Therefore the solver will try to find a displacement  $\Delta\hat{x}$ , such that  $C_i(\hat{x} + \Delta\hat{x}) = 0$ . A sensible choice is to move the particles a fraction  $\lambda$  along the constraint gradient  $\nabla C$  proportional to inverse inertia  $M^{-1}$ ;

$$\Delta\hat{x} = M^{-1} \nabla C \lambda. \quad (6)$$

An approximate solution to the linear equation above is derived using a method similar to a Jacobi iteration. The velocities are derived implicitly from the differences of the positions between frames. In the application PBF is used to replace the computation of the pressure term in the Navier-Stokes equation. In order to avoid attractive forces, which may lead to particle clumping and unwanted behaviour, the constraint from (5) is changed to an inequality constraint,  $C_i > 0$ .

## 2.3 Neighborhood Search

A lot of computation time is used for searching the neighborhood of a particle in both WCSPH and PBF. Efficient algorithms to find the particles in direct proximity is key to stable runtime performance. The application uses the CompactNSearch algorithm by Bender et al. [BK15], which uses the parallel, hash-based neighborhood search by Ihmsen et al. [IABT11]. To further enhance runtime performance, we perform a Z-index sort of particle positions after a fixed number of iterations. Sorting the particles depending on their spatial properties will decrease the number of cache misses and therefore speed up the access time.

## 2.4 Time integration

The system is advanced in time using the semi-implicit euler method. It is also known as the symplectic euler, as it conserves the energy of the system. XSPH smoothing is employed which aligns velocities of neighboring particles. This results in a viscosity effect similar to that of the viscosity due to the velocity divergence and helps to keep the simulation stable. The time step  $\Delta t$  is determined by the Courant-Friedrich-Levy (CFL) condition,  $\Delta t \leq \lambda \frac{2r}{\|v_{max}\|}$ , with particle radius  $r$  and maximum velocity of all particles  $v_{max}$ . The CFL condition will reduce the timestep dynamically, so that the particles will only move a fraction of their diameter during one timestep.

## 2.5 Boundary Interaction

To realise the interaction between the fluid and rigid boundaries, we use the boundary particle method by Akinici et al. [AIA<sup>+</sup>12]. By taking advantage of the particle based nature of the presented methods we are able to integrate contributions from boundary particles into the calculations of densities, forces and position corrections.

As boundaries in a scene are usually represented by triangular mesh models, we sample each triangle of a given input mesh. To this end, it is reasonable to dissolve a higher resolution mesh into a low resolution mesh which still contains the shape information and use the low resolution mesh for the simulation and keep the original high resolution mesh for rendering the final simulation. This will reduce the overall number of triangles and therefore the number of sampled particles. Figure 1 shows the mesh used for most simulations in this report.

To sample the particles on a triangle surface we perform a variation of the Pineda algorithm [Pin88] in the local coordinate system of the triangle surface. Given a triangle specified by the three vertex points  $a, b, c \in \mathbb{R}^3$ , we construct the orthonormal base vectors  $u, v, n \in \mathbb{R}^3$  defining the triangle surface. The basis transform that will yield the local triangle coordinates  $p \in \mathbb{R}^2$  for a point in global space  $x \in \mathbb{R}^3$  is given in



Figure 1: High resolution mesh for rendering (left) and low resolution mesh for simulation (right). The mesh on the left maintains the qualitative shape of the mesh.

equation 7.

$$p = \begin{pmatrix} u_x & u_v & u_z & u^T a \\ v_x & v_v & v_z & v^T a \\ 0 & 0 & 0 & 1 \end{pmatrix} \cdot x \quad (7)$$

Note that the projected point is in homogeneous two dimensional space, as we are only interested in sampling the surface of the triangle. Applying the pineda algorithm is now straight forward. We first enclose the triangle by a bounding box by determining the minimal and maximal coordinates in the projected space. We then go along the bounding box and check for each sample point if its inside the triangle by checking against its edges, as described by Pineda [Pin88]. We use a hexagonal sampling pattern to ensure maximum coverage of the surface. Figure 2 demonstrates the hexagonal sampling in the bounds of the rectangle enclosing the triangle in its local coordinate system. Sample points inside the triangle are shaded opaque blue and sample points outside are shaded transparently. Note that sample points that slightly touch the border are also taken to be in the surface of the triangle. This is achieved by allowing an offset of the particle radius and is used to ensure that the triangle mesh of the boundary has no holes where the edges of two triangles meet. The oversampling is later compensated by recomputing the volume of each sample depending on the number of samples in its close neighbourhood.

Given the sample points computed in the




Figure 2: Hexagonal sampling using pineda algorithm in local triangle space.

two dimensional surface plane, we transform the points back into three dimensional global space by multiplying them by the inverse of the matrix given in equation 7. Since the matrix is orthonormal its inverse is given by simply transposing the matrix.

The boundary particle radius should optimally be the same as the fluid particle radius. This is due to the timestep being determined by the CFL condition. If the boundary particle radius is much smaller. The fluid particles may tunnel through the boundary if the timestep is too high. If the boundary particle radius is too high, the boundary shape information is lost and the interaction therefore incorrect. The constant factor  $\lambda$  in the CFL is set around 0.4 in all further simulations.

## 2.6 Tension and Adhesion

Surface tension effects are implemented after the method by Akinci et al. [AAT13]. Akinci et al. define two types of forces, namely fluid cohesion and boundary adhesion. Fluid cohesion is an attractive force between particles. Macroscopically, it acts as a force field around the surface of the fluid which keeps the particles together, causing surface tension. The adhesion force causes fluid particles to stick to solid boundaries, which results in a wetting effect.

## 2.7 Surface Reconstruction

To extract an explicit representation of the fluid surface we first have to construct a signed distance function (SDF)  $\Phi(\vec{x})$  describing the isosurface. We use the SPH method to approximate the isosurface at any arbitrary point  $\vec{x}$ , depending on the normalized density  $\rho_j$  of the neighborhood region  $j \in N$ ;

$$\Phi(\vec{x}) = \sum_{j \in N(\vec{x})} \frac{1}{\rho_j} W(\|\vec{x} - x_j\|/h) - c. \quad (8)$$

The parameter  $c$  controls the zero level of the isosurface and therefore the enclosed volume of the explicit representation. We set  $c = 0.6$  during the remainder of the report. Using this method, we compute  $\Phi(\vec{x})$  at points in grid and reconstruct the triangle mesh representing the surface using the Marching Cubes algorithm [LC87].

## 2.8 Implementation

The implementation uses a C-style imperative approach with the exception of a few object oriented paradigms. Data encapsulation in the form of C++ classes is applied to ensure data consistency during computation. Basic inheritance is used to be able to reuse code for different parts of the program. This style of programming allows us to be flexible enough to incrementally implement the practical assignment tasks while sacrificing little to no performance due to design overload.

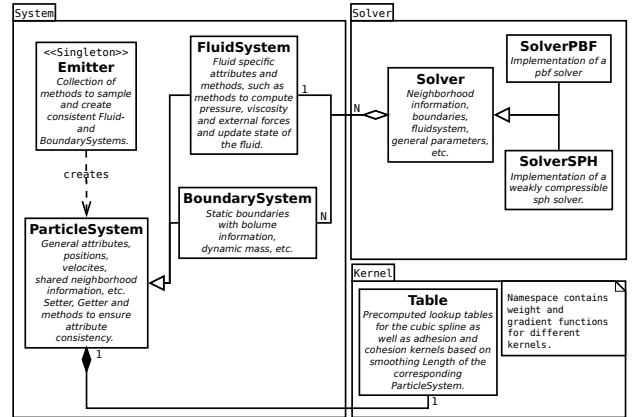


Figure 3: Schematic overview of the implementation.

The implementation is divided into two broad categories, systems and solvers. What we define as a 'system', includes fluid- and boundary particles, their generalization as particlesystems and auxiliary constructs concerning particles. These represent the creation, state and behaviour of fluid- and boundary particles. On the other hand, we have two separate specific solvers for WCSPH fluids and PBF. They aggregate fluid- and particlesystems to model the progression of the given system. Figure 3 gives a schematic UML diagram of the relation between the different parts of the implementation.

A particlesystem is a collection of particles which are represented by basic attributes, such as mass, density, position, velocity and the accumulated external forces on the particles. Since it is vital to keep all attributes of a particlesystem consistent during computation a particlesystem object is only allowed to be created by the emitter object. The emitter object is a singleton, which means that there is always one and only one single instance of the object during the runtime of the program. During runtime a particlesystem is present in two forms, either as fluid particles or as boundary particles. These subtypes possess attributes and methods responsible for the behaviour of fluid or boundary particles in a system. For example, boundary particles possess a separate variable volume attribute and a method to compute the volume for each particle depending on the density, e.g. the number of boundary particles in close neighbourhood.

## Results

## Conclusion

## Outlook

The WCSPH and PBF implementation as well as the surface reconstruction offers a high level of concurrency in the form of data parallelism. This means, that there are many parts in the program where sequentially processed data can be split into independent groups which can be processed in parallel.

## References

- [AAT13] Nadir Akinci, Gizem Akinci, and Matthias Teschner. Versatile surface tension and adhesion for sph fluids. *ACM Transactions on Graphics*, 10 2013.
- [AIA<sup>+</sup>12] Nadir Akinci, Markus Ihmsen, Gizem Akinci, Barbara Solenthaler, and Matthias Teschner. Versatile rigid-fluid coupling for incompressible sph. *ACM Trans. Graph.*, 31:62:1–62:8, 07 2012.
- [BK15] Jan Bender and Dan Koschier. Divergence-free smoothed particle hydrodynamics. In *Proceedings of the 14th ACM SIGGRAPH/Eurographics symposium on computer animation*, pages 147–155, 2015.
- [IABT11] Markus Ihmsen, Nadir Akinci, Markus Becker, and Matthias Teschner. A parallel sph implementation on multi-core cpus. In *Computer Graphics Forum*, volume 30, pages 99–112. Wiley Online Library, 2011.
- [IOS<sup>+</sup>14] Markus Ihmsen, Jens Orthmann, Barbara Solenthaler, Andreas Kolb, and Matthias Teschner. SPH Fluids in Computer Graphics. In Sylvain Lefebvre and Michela Spagnuolo, editors, *Eurographics 2014 - State of the Art Reports*. The Eurographics Association, 2014.
- [LC87] William E Lorensen and Harvey E Cline. Marching cubes: A high resolution 3d surface construction algorithm. *ACM siggraph computer graphics*, 21(4):163–169, 1987.
- [MHHR07] Matthias Müller, Bruno Heidelberger, Marcus Hennix, and John Ratcliff. Position based dynamics. *Journal of Visual Communication and Image Representation*, 18(2):109–118, 2007.
- [MM13] Miles Macklin and Matthias Müller. Position based fluids. *ACM Transactions on Graphics (TOG)*, 32(4):1–12, 2013.
- [Mon92] Joe J Monaghan. Smoothed particle hydrodynamics. *Annual review of astronomy and astrophysics*, 30(1):543–574, 1992.
- [Pin88] Juan Pineda. A parallel algorithm for polygon rasterization. In *Proceedings of the 15th annual conference on Computer graphics and interactive techniques*, pages 17–20, 1988.